

Optimizing Early Diabetes Detection: A Comparative Analysis of Traditional Machine Learning Baseline and Enhanced Deep Learning Models

###Diabetes Detection Tool

Author: Muhammad Waqas Mandrey

Section 1: Explore Data

In this section, we explore the data and check for any missing values in the dataset that may need to be addressed and any features consisting of values that may pose redundancy.

```
# Importing libraries
import pandas as pd
import numpy as np
import tensorflow as tf

np.random.seed(42)
tf.random.set_seed(42)

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

!ls /content/drive/MyDrive/DS402MLProj/

402MLProj.ipynb  diabetesDataset.csv

data =
pd.read_csv('/content/drive/MyDrive/DS402MLProj/diabetesDataset.csv')
data.iloc[500:]
```

	Age	Gender	Polyuria	Polydipsia	sudden weight loss	weakness
Polyphagia \						
500	66	Male	Yes	No	Yes	No
No						
501	67	Male	No	No	No	No
Yes						
502	70	Male	Yes	No	No	No
Yes						
503	44	Male	No	No	No	No
No						
504	38	Male	No	No	No	No
No						
505	35	Male	No	No	No	No
No						
506	61	Male	No	No	No	Yes

No							
507	60	Male	No	No	Yes	No	
No							
508	58	Male	No	No	No	Yes	
No							
509	54	Male	No	No	No	No	
No							
510	67	Male	No	No	No	Yes	
No							
511	66	Male	No	No	No	Yes	
Yes							
512	43	Male	No	No	No	No	
No							
513	62	Female	Yes	Yes	Yes	Yes	
No							
514	54	Female	Yes	Yes	Yes	Yes	
Yes							
515	39	Female	Yes	Yes	Yes	No	
Yes							
516	48	Female	Yes	Yes	Yes	Yes	
Yes							
517	58	Female	Yes	Yes	Yes	Yes	
Yes							
518	32	Female	No	No	No	Yes	
No							
519	42	Male	No	No	No	No	
No							

Genital thrush visual blurring Itching Irritability delayed healing \				
500	Yes	No	Yes	Yes
No				
501	No	Yes	No	No
No				
502	No	Yes	Yes	No
Yes				
503	No	No	No	No
No				
504	No	No	No	No
No				
505	No	No	No	No
No				
506	Yes	No	Yes	No
Yes				
507	No	No	No	No
No				
508	No	No	Yes	No
Yes				
509	No	No	No	No

No				
510	No	No	Yes	No
Yes				
511	No	Yes	Yes	No
Yes				
512	No	No	No	No
No				
513	No	Yes	No	No
No				
514	No	No	No	No
No				
515	No	No	Yes	No
Yes				
516	No	No	Yes	Yes
Yes				
517	No	Yes	No	No
No				
518	No	Yes	Yes	No
Yes				
519	No	No	No	No
No				

	partial paresis	muscle stiffness	Alopecia	Obesity	class
500	No	No	Yes	No	Positive
501	No	Yes	No	No	Negative
502	Yes	Yes	Yes	No	Negative
503	No	No	No	No	Negative
504	No	No	No	No	Negative
505	No	No	No	No	Negative
506	No	No	Yes	No	Negative
507	No	No	No	Yes	Negative
508	No	Yes	Yes	No	Negative
509	No	No	No	No	Negative
510	No	No	Yes	No	Negative
511	Yes	Yes	Yes	No	Negative
512	No	No	Yes	No	Negative
513	Yes	No	No	Yes	Positive
514	Yes	No	No	No	Positive
515	Yes	No	No	No	Positive
516	Yes	No	No	No	Positive
517	Yes	Yes	No	Yes	Positive
518	No	No	Yes	No	Negative
519	No	No	No	No	Negative

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 520 entries, 0 to 519
```

```
Data columns (total 17 columns):
```

```
#    Column                Non-Null Count  Dtype
```

```

---
0   Age                520 non-null    int64
1   Gender              520 non-null    object
2   Polyuria            520 non-null    object
3   Polydipsia          520 non-null    object
4   sudden weight loss  520 non-null    object
5   weakness            520 non-null    object
6   Polyphagia          520 non-null    object
7   Genital thrush      520 non-null    object
8   visual blurring     520 non-null    object
9   Itching             520 non-null    object
10  Irritability        520 non-null    object
11  delayed healing     520 non-null    object
12  partial paresis     520 non-null    object
13  muscle stiffness    520 non-null    object
14  Alopecia            520 non-null    object
15  Obesity             520 non-null    object
16  class               520 non-null    object
dtypes: int64(1), object(16)
memory usage: 69.2+ KB

data.columns

Index(['Age', 'Gender', 'Polyuria', 'Polydipsia', 'sudden weight
loss',
      'weakness', 'Polyphagia', 'Genital thrush', 'visual blurring',
      'Itching', 'Irritability', 'delayed healing', 'partial
paresis',
      'muscle stiffness', 'Alopecia', 'Obesity', 'class'],
      dtype='object')

# Number of columns in the dataframe
len(data.columns)

17

# Here we print our columns
print(data.iloc[:, 1:6])
print(data.iloc[:, 6:11])
print(data.iloc[:, 11:16])
print(data.iloc[:, 16:])

   Gender Polyuria Polydipsia sudden weight loss weakness
0    Male      No      Yes          No      Yes
1    Male      No      No          No      Yes
2    Male     Yes      No          No      Yes
3    Male      No      No          Yes      Yes
4    Male     Yes     Yes          Yes      Yes
..    ...     ...     ...     ...     ...
515  Female     Yes     Yes          Yes      No
516  Female     Yes     Yes          Yes      Yes

```

517	Female	Yes	Yes	Yes	Yes
518	Female	No	No	No	Yes
519	Male	No	No	No	No

[520 rows x 5 columns]

	Polyphagia	Genital thrush	visual blurring	Itching	Irritability
0	No	No	No	Yes	No
1	No	No	Yes	No	No
2	Yes	No	No	Yes	No
3	Yes	Yes	No	Yes	No
4	Yes	No	Yes	Yes	Yes
..
515	Yes	No	No	Yes	No
516	Yes	No	No	Yes	Yes
517	Yes	No	Yes	No	No
518	No	No	Yes	Yes	No
519	No	No	No	No	No

[520 rows x 5 columns]

	delayed healing	partial paresis	muscle stiffness	Alopecia	Obesity
0	Yes	No	Yes	Yes	Yes
1	No	Yes	No	Yes	No
2	Yes	No	Yes	Yes	No
3	Yes	No	No	No	No
4	Yes	Yes	Yes	Yes	Yes
..
515	Yes	Yes	No	No	No
516	Yes	Yes	No	No	No
517	No	Yes	Yes	No	Yes
518	Yes	No	No	Yes	No
519	No	No	No	No	No

[520 rows x 5 columns]

	class
0	Positive
1	Positive
2	Positive
3	Positive
4	Positive
..	...
515	Positive
516	Positive
517	Positive
518	Negative
519	Negative

[520 rows x 1 columns]

identifying columns with binary values

cols = ['Gender', 'Polyuria', 'Polydipsia', 'sudden weight loss',

```

        'weakness', 'Polyphagia', 'Genital thrush', 'visual blurring',
        'Itching', 'Irritability', 'delayed healing', 'partial
 paresis',
        'muscle stiffness', 'Alopecia', 'Obesity', 'class']

for i in cols:
    unique_vals = data[i].unique()
    print(unique_vals)

['Male' 'Female']
['No' 'Yes']
['Yes' 'No']
['No' 'Yes']
['Yes' 'No']
['No' 'Yes']
['No' 'Yes']
['No' 'Yes']
['Yes' 'No']
['No' 'Yes']
['Yes' 'No']
['No' 'Yes']
['Yes' 'No']
['Yes' 'No']
['Yes' 'No']
['Positive' 'Negative']

```

Section 2: Feature Engineering

In this section we perform feature engineering on our dataset. Here we begin with performing one-hot encoding to convert our binary categorical features into numerical features.

```

cat_col = ['Gender', 'Polyuria', 'Polydipsia', 'sudden weight loss',
           'weakness', 'Polyphagia', 'Genital thrush', 'visual blurring',
           'Itching', 'Irritability', 'delayed healing', 'partial
 paresis',
           'muscle stiffness', 'Alopecia', 'Obesity', 'class']

```

```

data = pd.get_dummies(data, columns = cat_col, drop_first = True)
data

```

	Age	Gender_Male	Polyuria_Yes	Polydipsia_Yes	sudden weight loss_Yes
0	40	1	0	1	
1	58	1	0	0	
2	41	1	1	0	
3	45	1	0	0	

1				
4	60	1	1	1
1				
..
...				
515	39	0	1	1
1				
516	48	0	1	1
1				
517	58	0	1	1
1				
518	32	0	0	0
0				
519	42	1	0	0
0				

	weakness_Yes	Polyphagia_Yes	Genital thrush_Yes	visual
blurring_Yes \				
0	1	0	0	
0				
1	1	0	0	
1				
2	1	1	0	
0				
3	1	1	1	
0				
4	1	1	0	
1				
..	
...				
515	0	1	0	
0				
516	1	1	0	
0				
517	1	1	0	
1				
518	1	0	0	
1				
519	0	0	0	
0				

	Itching_Yes	Irritability_Yes	delayed healing_Yes	partial
paresis_Yes \				
0	1	0	1	
0				
1	0	0	0	
1				
2	1	0	1	
0				

3	1	0	1
0			
4	1	1	1
1			
..
...			
515	1	0	1
1			
516	1	1	1
1			
517	0	0	0
1			
518	1	0	1
0			
519	0	0	0
0			

	muscle stiffness_Yes	Alopecia_Yes	Obesity_Yes	class_Positive
0	1	1	1	1
1	0	1	0	1
2	1	1	0	1
3	0	0	0	1
4	1	1	1	1
..
515	0	0	0	1
516	0	0	0	1
517	1	0	1	1
518	0	1	0	0
519	0	0	0	0

[520 rows x 17 columns]

data.columns

```
Index(['Age', 'Gender_Male', 'Polyuria_Yes', 'Polydipsia_Yes',
      'sudden weight loss_Yes', 'weakness_Yes', 'Polyphagia_Yes',
      'Genital thrush_Yes', 'visual blurring_Yes', 'Itching_Yes',
      'Irritability_Yes', 'delayed healing_Yes', 'partial
      paresis_Yes',
      'muscle stiffness_Yes', 'Alopecia_Yes', 'Obesity_Yes',
      'class_Positive'],
      dtype='object')
```

here we shuffle our data to make it more meaningful

```
data = data.sample(n = len(data))
data = data.reset_index(drop = True)
data
```

Age	Gender_Male	Polyuria_Yes	Polydipsia_Yes	sudden weight loss_Yes \
-----	-------------	--------------	----------------	-----------------------------

0	72	1	1	0
0				
1	40	0	1	1
0				
2	57	1	1	1
0				
3	41	1	1	1
1				
4	45	0	0	0
0				
..
...				
515	35	0	0	1
1				
516	58	1	0	1
1				
517	40	0	1	1
1				
518	57	1	1	1
1				
519	90	0	0	1
1				

	weakness_Yes	Polyphagia_Yes	Genital thrush_Yes	visual
blurring_Yes \				
0	0	1	0	
1				
1	1	1	0	
0				
2	1	1	1	
0				
3	1	1	1	
1				
4	0	1	0	
1				
..	
...				
515	1	0	0	
0				
516	1	1	0	
1				
517	1	0	0	
1				
518	1	1	0	
1				
519	0	0	1	
1				

Itching_Yes Irritability_Yes delayed healing_Yes partial

```

paresis_Yes \
0      1      0      1
1
1      1      0      0
1
2      0      0      1
1
3      1      1      0
0
4      1      0      0
1
..      ...      ...      ...
...
515     1      0      1
1
516     1      0      0
1
517     0      0      1
1
518     0      0      0
1
519     1      0      0
0

      muscle_stiffness_Yes  Alopecia_Yes  Obesity_Yes  class_Positive
0              1              1              0              0
1              0              0              0              1
2              0              0              0              1
3              0              0              1              1
4              0              0              0              1
..              ...              ...              ...              ...
515             1              0              0              1
516             0              1              1              1
517             1              0              0              1
518             0              0              0              1
519             1              1              0              1

[520 rows x 17 columns]

```

Engineering Features Summary

```

cols_num = data.select_dtypes(include=['number']).shape[1]
cols_cat = data.select_dtypes(exclude=['number']).shape[1]

print('Total number of features:', len(data.columns))
print('Numerical Features:', cols_num)
print('Categorical Features:', cols_cat)

```

```
Total number of features: 17
Numerical Features: 17
Categorical Features: 0
```

Section 3: Splitting Data for Training, Validation & Testing

In this section we split our data into 70% training, 15% validation & 15% test data by fractions.

```
# here we store 30% of the data as validation and test data combined
df_val_test = data.sample(frac = 0.30)

# here we split the data_val_test into 50-50 valid and test data
df_val = df_val_test.sample(frac = 0.5)
df_test = df_val_test.drop(df_val.index)

# here we use the remaining data 50% as the training data
df_train = data.drop(df_val_test.index)

df_train.shape
(364, 17)

df_val.shape
(78, 17)

df_test.shape
(78, 17)
```

Here we preprocess our data as per training the baseline and further for the deep learning model

```
X_train = df_train.drop('class_Positive', axis = 1)
y_train = df_train['class_Positive']

X_val = df_val.drop('class_Positive', axis = 1)
y_val = df_val['class_Positive']

X_test = df_test.drop('class_Positive', axis = 1)
y_test = df_test['class_Positive']

# this code describes the shape of the training, validation, testing
sets
print('Training shapes:',X_train.shape, y_train.shape)
print('Validation shapes:',X_val.shape, y_val.shape)
print('Test shapes:',X_test.shape, y_test.shape)

Training shapes: (364, 16) (364,)
Validation shapes: (78, 16) (78,)
Test shapes: (78, 16) (78,)
```

Section 4: Feature Scaling & Baseline Model: Logistic Regression

In this section we conduct our initial prediction analysis implementing logistic regression as the baseline model.

```
# Here we import our libraries for baseline and further processing of
the data
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression

# Here we perform feature scaling as some baseline models, in our case
logistic regression, are sensitive to scale of input features
scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)

# Now we train our baseline model with the training data
bmodel = LogisticRegression(random_state = 42)

bmodel.fit(X_train, y_train)

LogisticRegression(random_state=42)

# Now we use the trained model to perform predictions on our
validation set
y_val_pred = bmodel.predict(X_val)

# Now we evaluate our model using the following functions
from sklearn.metrics import roc_auc_score, accuracy_score,
precision_score, recall_score
def calc_specificity(y_actual, y_pred, thresh):
    # calculates specificity
    return sum((y_pred < thresh) & (y_actual == 0)) / sum(y_actual == 0)

def print_report(y_actual, y_pred, thresh):

    auc = roc_auc_score(y_actual, y_pred)
    accuracy = accuracy_score(y_actual, (y_pred > thresh))
    recall = recall_score(y_actual, (y_pred > thresh))
    precision = precision_score(y_actual, (y_pred > thresh))
    specificity = calc_specificity(y_actual, y_pred, thresh)
    print('AUC: %.3f' % auc)
    print('accuracy: %.3f' % accuracy)
    print('recall: %.3f' % recall)
    print('precision: %.3f' % precision)
    print('specificity: %.3f' % specificity)
```

```

    print(' ')
    return auc, accuracy, recall, precision, specificity

# Here we set our threshold and call the model evaluation functions to
# determine the performance of the model with the validation set
thresh = 0.5
y_train_preds = bmodel.predict_proba(X_train)[: ,1]
y_val_preds = bmodel.predict_proba(X_val)[: ,1]

print('Logistic Regression')
print('Training:')
bmodel_train_auc, bmodel_train_accuracy, bmodel_train_recall, \
    bmodel_train_precision, bmodel_train_specificity =
print_report(y_train, y_train_preds, thresh)
print('Validation:')
bmodel_valid_auc, bmodel_valid_accuracy, bmodel_valid_recall, \
    bmodel_valid_precision, bmodel_valid_specificity =
print_report(y_val ,y_val_preds, thresh)

Logistic Regression
Training:
AUC:0.986
accuracy:0.945
recall:0.963
precision:0.946
specificity:0.917

Validation:
AUC:0.943
accuracy:0.897
recall:0.896
precision:0.935
specificity:0.900

# Now we use the trained model to perform predictions on our test set
y_test_pred = bmodel.predict(X_test)

# Here we call the model evaluation functions to determine the
# performance of the model with the test set
y_test_preds = bmodel.predict_proba(X_test)[: ,1]

print('Logistic Regression')
print('Test:')
bmodel_valid_auc, bmodel_valid_accuracy, bmodel_valid_recall, \
    bmodel_valid_precision, bmodel_valid_specificity =
print_report(y_test, y_test_preds, thresh)

Logistic Regression
Test:
AUC:0.966

```

```

accuracy:0.923
recall:0.925
precision:0.961
specificity:0.920

features = ['Age', 'Gender_Male', 'Polyuria_Yes', 'Polydipsia_Yes',
            'sudden weight loss_Yes', 'weakness_Yes', 'Polyphagia_Yes',
            'Genital thrush_Yes', 'visual blurring_Yes', 'Itching_Yes',
            'Irritability_Yes', 'delayed healing_Yes', 'partial
pareisis_Yes',
            'muscle stiffness_Yes', 'Alopecia_Yes', 'Obesity_Yes']

# Here we perform feature importance over our regression model to
identify the priority features for our model
feature_importances = pd.DataFrame(bmodel.coef_[0],
                                   index = features,

columns=['importance']).sort_values('importance',

ascending=False)

feature_importances.head()

```

	importance
Polydipsia_Yes	1.871165
Polyuria_Yes	1.623144
Irritability_Yes	1.152463
Genital thrush_Yes	0.932349
partial paresis_Yes	0.706734

Section 5: Design & Implement Deep Learning Model: Feed-Forward Neural Network

```

from tensorflow import keras

model = keras.Sequential([
    keras.layers.Input(shape=(16,)), # Input layer with 16 features
    keras.layers.Dense(50, activation='relu'), # We implement the
model with 2 hidden layers, with 50 neurons and 30 neurons
respectively
    keras.layers.Dense(30, activation='relu'), # We implement the
model with ReLU activation
    keras.layers.Dense(1, activation='sigmoid') # Output layer with 1
neuron and sigmoid activation for binary classification
])

# Here we compile our model with the optimizer and further parameters
model.compile(optimizer='Adam', # We use Adam optimizer
              loss='binary_crossentropy', # For our binary

```

```
classification data we implement with Binary cross-entropy loss
        metrics=['accuracy']) # Here we track accuracy of each
run
```

```
# Here we train our Feed-forward Neural Network Model
```

```
model.fit(X_train, y_train, epochs=15, validation_data=(X_val, y_val))
```

```
Epoch 1/15
```

```
12/12 [=====] - 2s 33ms/step - loss: 0.6530 -
accuracy: 0.6703 - val_loss: 0.6034 - val_accuracy: 0.7949
```

```
Epoch 2/15
```

```
12/12 [=====] - 0s 7ms/step - loss: 0.5281 -
accuracy: 0.8846 - val_loss: 0.5105 - val_accuracy: 0.8205
```

```
Epoch 3/15
```

```
12/12 [=====] - 0s 8ms/step - loss: 0.4419 -
accuracy: 0.9093 - val_loss: 0.4443 - val_accuracy: 0.8205
```

```
Epoch 4/15
```

```
12/12 [=====] - 0s 11ms/step - loss: 0.3723 -
accuracy: 0.9258 - val_loss: 0.3892 - val_accuracy: 0.8974
```

```
Epoch 5/15
```

```
12/12 [=====] - 0s 7ms/step - loss: 0.3155 -
accuracy: 0.9258 - val_loss: 0.3435 - val_accuracy: 0.9231
```

```
Epoch 6/15
```

```
12/12 [=====] - 0s 7ms/step - loss: 0.2711 -
accuracy: 0.9258 - val_loss: 0.3112 - val_accuracy: 0.9359
```

```
Epoch 7/15
```

```
12/12 [=====] - 0s 13ms/step - loss: 0.2373 -
accuracy: 0.9313 - val_loss: 0.2888 - val_accuracy: 0.9231
```

```
Epoch 8/15
```

```
12/12 [=====] - 0s 8ms/step - loss: 0.2101 -
accuracy: 0.9341 - val_loss: 0.2750 - val_accuracy: 0.9231
```

```
Epoch 9/15
```

```
12/12 [=====] - 0s 8ms/step - loss: 0.1895 -
accuracy: 0.9396 - val_loss: 0.2654 - val_accuracy: 0.9231
```

```
Epoch 10/15
```

```
12/12 [=====] - 0s 8ms/step - loss: 0.1714 -
accuracy: 0.9396 - val_loss: 0.2581 - val_accuracy: 0.9231
```

```
Epoch 11/15
```

```
12/12 [=====] - 0s 7ms/step - loss: 0.1563 -
accuracy: 0.9396 - val_loss: 0.2555 - val_accuracy: 0.8718
```

```
Epoch 12/15
```

```
12/12 [=====] - 0s 8ms/step - loss: 0.1436 -
accuracy: 0.9505 - val_loss: 0.2541 - val_accuracy: 0.8718
```

```
Epoch 13/15
```

```
12/12 [=====] - 0s 8ms/step - loss: 0.1324 -
accuracy: 0.9588 - val_loss: 0.2542 - val_accuracy: 0.8718
```

```
Epoch 14/15
```

```
12/12 [=====] - 0s 9ms/step - loss: 0.1223 -
accuracy: 0.9615 - val_loss: 0.2516 - val_accuracy: 0.8590
```

```
Epoch 15/15
```

```
12/12 [=====] - 0s 8ms/step - loss: 0.1144 - accuracy: 0.9615 - val_loss: 0.2508 - val_accuracy: 0.8590
```

```
<keras.src.callbacks.History at 0x7fd9f0cb4a30>
```

```
# Here we perform predictions on our test data with the trained model  
y_pred = model.predict(X_test)
```

```
3/3 [=====] - 0s 4ms/step
```

```
# Evaluate the model on the test set  
test_loss, test_accuracy = model.evaluate(X_test, y_test)  
print(f"Test Loss: {test_loss}, Test Accuracy: {test_accuracy}")
```

```
3/3 [=====] - 0s 5ms/step - loss: 0.1519 - accuracy: 0.9231  
Test Loss: 0.15193887054920197, Test Accuracy: 0.9230769276618958
```

Section 6: Model Evaluation Analysis

In this section we conduct analysis on our model to determine the accurate results of the model using analysis techniques including F-1 Score, Recall, Precision, AUC score, and AUC-ROC curve.

```
from sklearn.metrics import precision_score, recall_score, f1_score,  
accuracy_score, confusion_matrix, roc_auc_score
```

```
y_train_preds_nn = model.predict(X_train)  
y_val_preds_nn = model.predict(X_val)  
y_test_preds_nn = model.predict(X_test)
```

```
y_train_preds_binary = (y_train_preds_nn > 0.5).astype(int)  
y_val_preds_binary = (y_val_preds_nn > 0.5).astype(int)  
y_test_preds_binary = (y_test_preds_nn > 0.5).astype(int)
```

```
# Here we determine precision, recall, f1-score, AUC, and accuracy  
based on the training data performance of the model
```

```
precision_train = precision_score(y_train, y_train_preds_binary)  
recall_train = recall_score(y_train, y_train_preds_binary)  
f1_train = f1_score(y_train, y_train_preds_binary)  
accuracy_train = accuracy_score(y_train, y_train_preds_binary)  
auc_train = roc_auc_score(y_train, y_train_preds_nn)
```

```
# Here we determine precision, recall, f1-score, AUC, and accuracy  
based on the validation data performance of the model
```

```
precision_val = precision_score(y_val, y_val_preds_binary)  
recall_val = recall_score(y_val, y_val_preds_binary)  
f1_val = f1_score(y_val, y_val_preds_binary)  
accuracy_val = accuracy_score(y_val, y_val_preds_binary)  
auc_val = roc_auc_score(y_val, y_val_preds_nn)
```



```

# Here we determine precision, recall, f1-score, AUC, and accuracy
based on the testing data performance of the model
precision_test = precision_score(y_test, y_test_preds_binary)
recall_test = recall_score(y_test, y_test_preds_binary)
f1_test = f1_score(y_test, y_test_preds_binary)
accuracy_test = accuracy_score(y_test, y_test_preds_binary)
auc_test = roc_auc_score(y_test, y_test_preds_nn)

print('Training Set Metrics:')
print(f'Precision: {precision_train:.4f}')
print(f'Recall: {recall_train:.4f}')
print(f'F1-score: {f1_train:.4f}')
print(f'Accuracy: {accuracy_train:.4f}')
print(f'AUC: {auc_train:.4f}')
print('Confusion Matrix:')
print(confusion_matrix(y_train, y_train_preds_binary))

print('\nValidation Set Metrics:')
print(f'Precision: {precision_val:.4f}')
print(f'Recall: {recall_val:.4f}')
print(f'F1-score: {f1_val:.4f}')
print(f'Accuracy: {accuracy_val:.4f}')
print(f'AUC: {auc_val:.4f}')
print('Confusion Matrix:')
print(confusion_matrix(y_val, y_val_preds_binary))

print('\nTest Set Metrics:')
print(f'Precision: {precision_test:.4f}')
print(f'Recall: {recall_test:.4f}')
print(f'F1-score: {f1_test:.4f}')
print(f'Accuracy: {accuracy_test:.4f}')
print(f'AUC: {auc_test:.4f}')
print('Confusion Matrix:')
print(confusion_matrix(y_test, y_test_preds_binary))

# # Additionally we compute the confusion matrix
# cm = confusion_matrix(y_test, y_test_preds_binary)
# print('Confusion Matrix:')
# print(cm)

12/12 [=====] - 0s 2ms/step
3/3 [=====] - 0s 3ms/step
3/3 [=====] - 0s 6ms/step
Training Set Metrics:
Precision: 0.9680
Recall: 0.9680
F1-score: 0.9680
Accuracy: 0.9615
AUC: 0.9939
Confusion Matrix:

```

```
[[138  7]
 [  7 212]]
```

Validation Set Metrics:

Precision: 0.9111

Recall: 0.8542

F1-score: 0.8817

Accuracy: 0.8590

AUC: 0.9535

Confusion Matrix:

```
[[26  4]
 [ 7 41]]
```

Test Set Metrics:

Precision: 0.9434

Recall: 0.9434

F1-score: 0.9434

Accuracy: 0.9231

AUC: 0.9864


Confusion Matrix:

```
[[22  3]
 [ 3 50]]
```

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)
```

```
plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()
```



```
from sklearn.metrics import roc_curve

y_val_pred = model.predict(X_val)


fpr_val, tpr_val, thresholds_val = roc_curve(y_val, y_val_pred)
auc_val = roc_auc_score(y_val, y_val_pred)

fpr_test, tpr_test, thresholds_test = roc_curve(y_test, y_pred)
auc_test = roc_auc_score(y_test, y_pred)

print('Validation AUC: %.2f' % auc_val)
print('Test AUC: %.2f' % auc_test)

plt.plot(fpr_val, tpr_val, 'r-', label = 'Validation AUC: %.2f' % auc_val)
plt.plot(fpr_test, tpr_test, 'b-', label = 'Test AUC: %.2f' % auc_test)
plt.plot([0,1],[0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```

3/3 [=====] - 0s 4ms/step
Validation AUC:0.95
Test AUC:0.99



```
from sklearn.metrics import roc_curve

y_train_pred = model.predict(X_train)

fpr_train, tpr_train, thresholds_train = roc_curve(y_train,
y_train_pred)
auc_train = roc_auc_score(y_train, y_train_pred)

fpr_val, tpr_val, thresholds_val = roc_curve(y_val, y_val_pred)
auc_val = roc_auc_score(y_val, y_val_pred)

fpr_test, tpr_test, thresholds_test = roc_curve(y_test, y_pred)
auc_test = roc_auc_score(y_test, y_pred)

print('Training AUC:%.2f'%auc_train)
print('Validation AUC:%.2f'%auc_val)
print('Test AUC:%.2f'%auc_test)

plt.plot(fpr_train, tpr_train, 'g-',label = 'Training AUC:
%.2f'%auc_train)
```

```
plt.plot(fpr_val, tpr_val, 'r-', label = 'Validation AUC: %.2f' % auc_val)
plt.plot(fpr_test, tpr_test, 'b-', label = 'Test AUC: %.2f' % auc_test)
plt.plot([0,1],[0,1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```

12/12 [=====] - 0s 3ms/step
 Training AUC:0.99
 Validation AUC:0.95
 Test AUC:0.99

